

KGen User's Guide for KGen ver. 0.7.2

CONTENTS

[0. Getting-started](#)

[1. Introduction](#)

[2. How to use](#)

[2.1 Installation](#)

[2.1.1 Requirements](#)

[2.2 Command-line syntax](#)

[2.2 User interface for specifying a kernel to be extracted](#)

[2.3 Supporting kernel generation for large-sized software](#)

[2.3.1 Namepath](#)

[2.3.2 "include" information](#)

[2.3.2.1 INI sections applicable to each source file](#)

[2.3.2.2 INI sections applicable to all source files](#)

[2.3.2.3 INI sections applicable to KGen operations](#)

[2.3.3 User-provided "exclude" information](#)

[2.3.3.1 namepath section](#)

[3. Kernel extraction examples](#)

[3.1 simple example](#)

[3.2 simple-MPI example](#)

[3.3 simple-region example](#)

[4. Command line options](#)

[5. Known Issues](#)

[6. Changes from KGen ver. 0.6.3](#)

[6.1 User Interface](#)

[6.2 Major Improvements](#)

0. Getting-started

Before using KGen, please make sure that following Linux commands work on your system.

- python (>=2.7 and <3.0)

- cpp
- make
- strace

Following shows how to extract a kernel from an example in KGen distribution.

```
>> cd example/simple;      # move to an example directory
>> vi src/Makefile;       # Modify FC if required
>> make;                   # extract a kernel
>> cd kernel;              # move to a kernel directory
>> make;                   # build and run a kernel
```

Verification and timing results of kernel execution should be displayed on screen if all of above commands complete successfully.

“Make” command for kernel extraction actually runs following KGen command. The argument is mandatory providing callsite location information. First three options of “--cmd-clean”, “--cmd-build”, and “--cmd-run” are also mandatory to clean & build & run target application. The detail of this KGen command is explained in extraction example sections in this document below.


```
${KGen} \
  ${SRC_FILE}:${CALLSITE} \
  --cmd-clean "cd ${SRC_DIR}; make clean" \
  --cmd-build "cd ${SRC_DIR}; make build" \
  --cmd-run "cd ${SRC_DIR}; make run" \
  --timing repeat=10 \
  --invocation 0:0:1,0:0:2,0:0:3 \
  --check tolerance=1.0D-14 \
  --rebuild all
```

1. Introduction

KGen is a Python tool that extracts Fortran statements from a larger software and combines the extracted statements as a stand-alone application, aka kernel. It also saves input and output data that are used in the kernel for execution and verification purpose.

General workflow for using KGen is shown in following figure. (A)To run KGen, user needs to provide four arguments: a “callsite location” that specifies which part to be extracted and three types of Linux commands to clean/build/run target production application. Once KGen completes with success, it will produce a set of kernel source files (B) and a set of input and

output data files (B). With the kernel files and data files, user can immediately compile (C) and run (D) the kernel as a stand-alone application and its outcome from execution is automatically verified against the generated output data.



In practice, user may adopt iterative workflow in using KGen. It is generally unclear which external library should be excluded or how much state data should be generated at first. User first provides KGen with minimum information and tries to generate a kernel. If it fails, in return, KGen may provide an user with what information was missed or what went wrong. In addition, KGen tries to create a kernel even though it can not find all required information and user may investigate the generated kernel further on missed information. User may re-try KGen by providing additional information, and repeat the steps.

This version of KGen is an integration of two KGen sub-applications running behind: Koption and Kext. Each KGen sub-application can be executed as a stand-alone too. Koption automates the generation of macro definitions and include paths. Kext extracts a kernel using information generated by Koption. If you have used previous version of KGen, Kext is basically same to KGen of previous version.

KGen is being developed actively as of writing this document. It would be very helpful if you can share your experience on using KGen by sending relevant information to "kgen@ucar.edu".

2. How to use

2.1 Installation

KGen is written in Python as an extension of a Fortran parser distributed in F2PY python package.

2.1.1 Requirements

KGen is developed mainly in Python 2.7. A Fortran parser of F2PY is extensively used in KGen and is a part of KGen distribution. Therefore, there is no additional work needed for installing the F2PY Fortran Parser. From this version of KGen, it relies on “strace” Linux utility. Other frequently used Linux utilities are “cpp” C preprocessor and “make” building utility.

Once you have unzipped KGen distribution file, then you can run KGen as below:

```
$PATH_TO_KGen/bin/kgen --help
```

If you are an user of Yellowstone cluster of NCAR, you can use a pre-installed KGen at “/glade/u/tdd/asap/contrib/kgen/std”

2.2 Command-line syntax

```
>> $PATH_TO_KGen/bin/kgen [KGen options] <filepath>[:<namepath>] --cmd-clean  
<commands> --cmd-build <commands> --cmd-run <commands>
```

[KGen options]: These are optional information. Please see a section below for details

<filepath>[:<namepath>]: This argument specify a file path that contains a callsite and/or namepath. Please see a section next for details.

--cmd-clean <commands>: This mandatory option is Linux command(s) that ensure that next “build commands” actually compiles all source files that KGen requires to analyze.

--cmd-build <commands>: This mandatory option is Linux command(s) that compiles target application.

--cmd-run <commands>: This mandatory option is Linux command(s) that executes target application.

If there are multiple Linux commands, use semicolon between them. Use quotation marks if there are spaces in commands.

example:

```
kgen callsite.F90 --cmd-clean "cd $WORK; make clean" --cmd-build "cd $WORK; make build"  
--cmd-run "cd $WORK; make run"
```

2.2 User interface for specifying a kernel to be extracted

KGen supports two ways of user-interface for specifying a kernel location to be extracted: command-line and KGen directive.

KGen directives are inserted in source codes to direct what KGen to do.

2.2.1 KGen directive user interface

In general, the syntax of KGen directive follows OpenMP syntax. Case is not sensitive if not specified.

SYNTAX: !\$kgen directive name/clause(s)

“directive” specifies the type of action. “name” or “clause(s)” is additional information for the directive. “name” is a non-space ascii string and “clause” is a name with having parenthesis.

Continuation of kgen directive is not supported yet.

Three KGen directives are implemented as of this version.

“callsite” directive

syntax: !\$kgen callsite name

meaning: callsite directive specifies the location of callsite Fortran statement in a source file. If this directive is specified in source file, user does not have to provide “namepath” on command-line. “namepath” is a colon-separated names. See “Namepath” section for more about “namepath”

The directive should be placed just before the callsite line. However, blank line(s) and other comment lines are allowed in-between. name is a user-provided string for kernel name.

example:

```
!$kgen callsite calc  
CALL calc(i, j, output)
```

“begin_callsite” and “end_callsite” directives

syntax:

```
!$kgen begin_callsite name
... fortran statements...
!$kgen end_callsite
```

meaning: begin_callsite and end_callsite directives specify a region of Fortran statements in a source file to be extracted as a kernel.

example:

```
!$kgen begin_callsite calc
DO i=1, COL
  DO j=1, ROW
    CALL calc(i, j, output)
  END DO
END DO
!$kgen end_callsite calc
```

“write” directive

syntax: !\$kgen write variable[,variable,...]

meaning: write directive specifies variables whose content will be saved in state data files and will be read by a generated kernel. By using this directive, user can manually force KGen to save state data at arbitrary points of source code. Multiple variables can be specified by using comma in-between. Variable can be a member of derived type. One restriction of specifying variable is that it can not be an element of array but array itself. For example if A%B is an array it can not be specified as A%B[index].

The directive can be located anywhere within executable part of source code.

example:

```
!$kgen write i,j
CALL calc(i, j, output)
```

“exclude” directive

syntax: !\$kgen exclude

meaning: exclude directive specifies that next statement will be excluded during kernel generation. The main purpose of this directive is to support MPI “receiving” routines such as “recv” within generated kernel together with “write” directive. For example, user can use “write” directive to force KGen to read/write a variable and “exclude” MPI receiving routine.

example:

```
!$kgen write data
!$kgen exclude
CALL MPI_RECV(data, count, MPI_DOUBLE_PRECISION, from, tag,
MPI_COMM_WORLD, status, ierr )
```

2.2.2 KGen command-line user interface

KGen can be invoked using command line interface too.

```
>> $PATH_TO_KGen/bin/kgen [KGen options] <filepath>:<namepath> <clean commands>
<build commands> <run commands>
```

“filepath” is a path to a source file that contains a code region to be extracted.” “namepath” tells KGen of the region of code to be extracted. Please see namepath section 2.3.1 below for more detail.

2.3 Supporting kernel generation for large-sized software

2.3.1 Namepath

To resolve possible name conflict among different levels of namespace, KGen introduces a hierarchical representation of a name, called “namepath.”

Namepath is consecutive names with colons between them as a separator. For example, if name B is inside of A, then it can be represented by “A:B”. In practice, it is used to specify “identifiers” in KGen. For example, a kernel extraction region can be specified as following:

```
example)
module A
    subroutine B
        CALL C(...)
    end subroutine B
end module A
```

To specify “CALL C(…)” statement in above sample code, user can use “A:B:C” namepath.

To increase the usability of “namepath”, several syntactic features are added to above basic usage.

The separator of colon can be used as a metacharacter meaning of “any” similar to “*” in “ls” linux command. First, leading colon means any names can be placed before a name placed next to the colon. For example, “:name_a” means any namepath that ends with “name_a”. Similarly, colon at the end of a namepath means any names can be followed after a name placed before the colon. For example, “name_a:” matches to any namepath that starts with “name_a”. Finally, double colons between names means any names can be placed between the two names. For example, “name_a::name_b” matched to any namepaths that starts with “name_a” and ends with “name_b”

Namepath examples)

C => A name that has only one-level whose name is “C” such as “module C”

:C => any name ends with “C” such as any variable in a subroutine in a module

C: => any name whose top-level name is “C” and may contains lower-level names such as all variables in a subroutine of “C”

:C: => any names of “C” in any levels

A::C => Any names whose top-level name is A and whose lowest-level name is C

2.3.2 “include” information

To analyze source code correctly, KGen requires to know what are macro definitions and “include” paths per each source file. KGen collects the information automatically through building target application under KGen control using “strace” utility. Once KGen collects the information, it generates “include.ini” text file in working directory. While, in simple case, user does not need to know the content of the file, there are cases that user-provided information in the file can help KGen to extract a kernel correctly and/or more efficiently.

Syntax of the INI file follows conventional INI file syntax. Brackets are used to specify sections. In a section, an option is added in a line or over multiple lines. Each option has a format of key and value pair with a separator of “=”. Value part can be missed depending on the type of option.

2.3.2.1 INI sections applicable to each source file

Some information has to be provided per each source file separately. As of this version, there are four types of information are identified in this category: macro definition, include directory, compiler path and compiler options. Following convention is used to provide these information in the INI file.

[Path-to-source-file]

include = [directory path]:[directory path]:...

macro_name = macro_value

...


```
compiler = path/to/compiler  
compiler_options = compiler options
```

example) When “program.F90” uses a module in “./module” directory, and “program.F90” needs macro definition of “N=10” with intel Fortran compiler and “-O2 -fp-model precise” option

```
[program.F90]  
include = ./module  
N = 10  
compiler = ifort  
compiler_options = -O2 -fp-model precise
```

There can be multiple macro_name options but only one “include”, “compiler”, and “compiler_options” option is allowed per each file.

2.3.2.2 INI sections applicable to all source files

There are several types of information that can be applied to all source files that KGen analyzes.

Common macro definitions and include directories

In some cases, all source files may share the same macro definitions and/or include directories. In the cases, instead of specifying the information per each source file sections, user can use following sections.

```
[macro]  
macro_name = macro_value  
...
```

These macros will be added to each source file during KGen analysis.

```
[include]  
include_path1 =  
include_path2 =  
...
```

These include_paths will be added to each source file during KGen analysis. Note that each path should be specified per each line, which is different from the syntax of separate section for each file. Value part of each option should be blank for this version and is reserved for later use.

2.3.2.3 INI sections applicable to KGen operations

User can set a compiler command and compiler options to be used in extracted kernel.

```
[compiler]
compiler = path/to/compiler
compiler_options = compiler options
```

User can provide additional information to import source files or object files or library files through “import” section in a INI file.

```
[import]
filepath = action
```

“source” and “object” and library actions are implemented as of this version. “source” action in import section provides KGen with paths to additional files to be analyzed before starting main parsing tasks. “object” action specifies a path to an object file that will be copied to “kernel” output directory. “Library” actions specifies a path to a folder contains libraries and name of library(similar to -L and -l compiler option)

```
[import]
/path/to/source/file.F90 = source
/path/to/object/file.o = object
/path/to/folder/contains/library_files = library(libname)
```

2.3.3 User-provided “exclude” information

KGen accepts an INI-format file with “-e” command-line option. In the INI file, user can provide KGen with information of names(or namepaths) to be excluded during name search. Details of using the INI file are explained in following sections.

Command line option format: “-e <user-providing INI format file>”

Syntax of the INI file follows conventional INI file syntax. Brackets are used to specify sections. In a section, a option is added in a line or over multiple lines.

2.3.3.1 namepath section

When “namepath” is specified in a section of INI file, “actions” specified in an option are applied to namepath in the option.

```
[namepath]
```

namepath = [action]

“namepath” in an option line specifies target of action. The syntax of “namepath” is explained in section 2.3.1.

Regardless of actions specified in an option line, any name in execution part of Fortran source codes that matches to namepaths will be skipped from name resolution in KGen. This is also true if there is no action is specified.

There are two “actions” defined in this version.

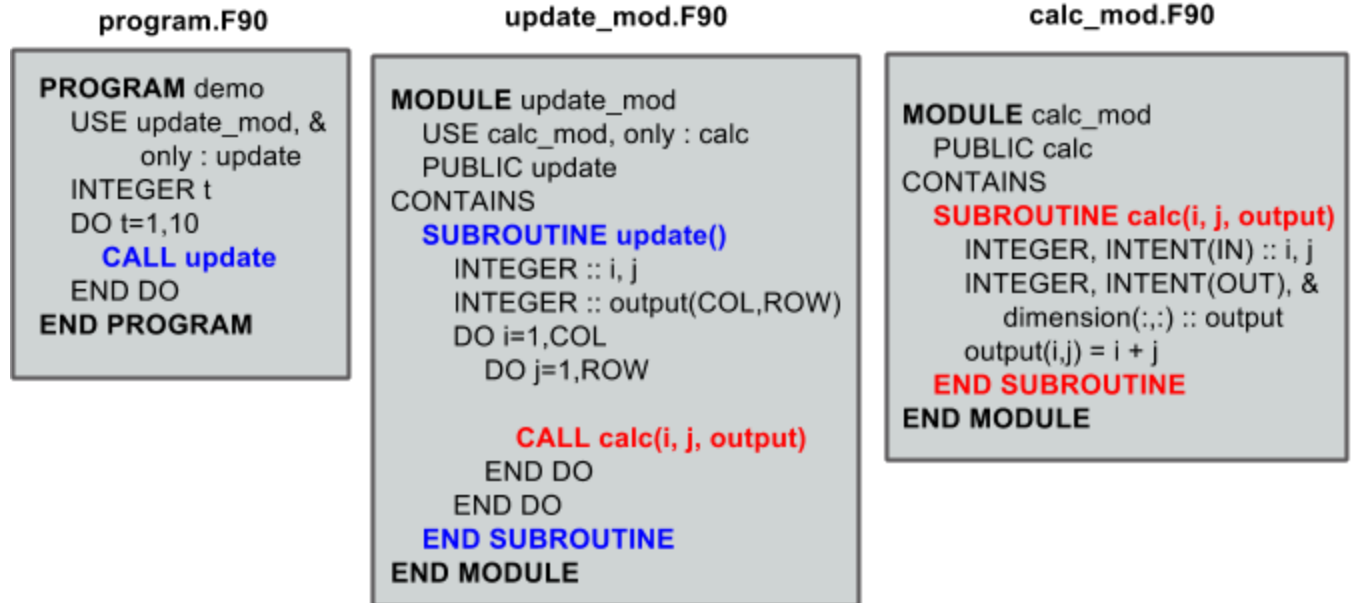
When “skip_module” is specified as “action”, KGen will not use a module specified by the namepath during name resolution tasks. This actions is useful when an module implemented in external library is used but not relevant to kernel extraction. By specifying this action, user can prune search tasks.

When “remove_state” is specified as “action”, KGen will not save state data specified by the namepath. This action may be useful discarding components of a derived type. If not all components of a derived type should be used for saving state data, user can specify components of the derived to be excluded from the state saving. This action may be used together with above “skip_module” action.

3. Kernel extraction examples

3.1 simple example

In this section, an example of generating a kernel of “calc” from “calc_mod.F90” source file is explained. This example shows a basic workflow of using KGen with information on using essential KGen options. You can find this example in “simple” folder of KGen distribution.



In this example, callsite is “CALL calc(i, j, output)” in update_mod.F90.

NOTE: Following macros are assumed in following command line.

SRC_DIR := \${PWD}/src

SRC := \${SRC_DIR}/update_mod.F90

CALLSITE := update_mod:update:calc

To extract “calc” subroutine and make it stand-alone application, following KGen command is used.

```

>> ${KGen} \
    ${SRC}:${CALLSITE} \
    --cmd-clean "cd ${SRC_DIR}; make clean" \
    --cmd-build "cd ${SRC_DIR}; make build" \
    --cmd-run "cd ${SRC_DIR}; make run" \
    --timing repeat=10 \
    --invocation 0:0:1,0:0:2,0:0:3 \
    --check tolerance=1.0D-14 \
    --rebuild all

```

Each lines of the command are explained below.

“>> \${KGen}” : This invokes KGen. It is assumed that \${KGen} has valid path to “kgen” in “bin” directory.

“\${SRC}/\${CALLSITE}” : This argument specifies the path to a source file that contains “callsite” and namepath to callsite.

`--cmd-clean "cd ${SRC_DIR}; make clean"`: This option is Linux commands that ensure that next build argument actually compiles all source files that KGen analyzes.

`--cmd-build "cd ${SRC_DIR}; make build"`: This option is Linux commands that compiles all source files to build target application. KGen collects macro definitions, include path, compiler and compiler options per each source files.

`--cmd-run "cd ${SRC_DIR}; make run"`: This option is Linux commands that runs target application. This command is used in KGen to generate state data files.

`--timing repeat=10`: `--timing` flag specifies information for timing measurement. `repeat` sub-flag specifies how many calls to be made to kernel to measure timing. Larger number increases the resolution of measurement but increase time to measure.

`--invocation 0:0:1,0:0:2,0:0:3`: `--invocation` flag specifies when to save state data in terms of calls to the kernel. In this example, state data will be saved first, second, and third calls to the kernel. First and second numbers in colon separated tripes represents MPI ranks and OpenMP number each. `0` is used for non-MPI application and non-OpenMP application as like this example.

`--check tolerance=1.0D-14`: This option sets the tolerance of normalized-RMS error check. If an error exceeds this value, KGen displays the test as failure. Without this option, KGen sets the value as `1.0D-15`.

`--rebuild all`: This option forces KGen to generate intermittent files such as strace log files and include.ini files. This option used for the purpose of demonstration in this example. Without this option, KGen reuses intermittent files. Please see a section below that explains all of KGen options for more sub-options.

Once KGen completes kernel generation successfully, kernel files and state data files are generated in `kernel` folder as shown below.

kernel folder

```
kernel
|-- calc.0.0.1
|-- calc.0.0.2
|-- calc.0.0.3
|-- calc_mod.F90
|-- kernel_driver.f90
|-- kgen_statefile.lst
|-- kgen_utils.f90
|-- Makefile
`-- update_mod.F90
```

Three data files, “calc.0.0.1”, “calc.0.0.2”, and “calc.0.0.3” in this example, are created in “kernel” folder. These state data files are used to drive kernel execution and to verify the result of kernel execution.

Final step is to build and run the “calc” kernel. Move to “kernel” folder and run “make”. If it completes with success, you will see something similar to following figure on screen.

```

>> make
ifort -c -o kgen_utils.o kgen_utils.f90
ifort -c -o calc_mod.o calc_mod.F90
ifort -c -o update_mod.o update_mod.F90
ifort -c -o kernel_driver.o kernel_driver.f90
ifort -o kernel.exe kernel_driver.o
      update_mod.o kgen_utils.o calc_mod.o

./kernel.exe

***** Verification against 'calc.0.0.1' *****

Number of verified variables:      3
Number of identical variables:     2
Number of non-identical variables within tolerance:  1
Number of non-identical variables out of tolerance:  0
Tolerance:  1.0000000000000000E-010

Verification PASSED

calc : Time per call (usec):  3.700000047683716E-002

***** Verification against 'calc.0.0.2' *****
...
***** Verification against 'calc.0.0.3' *****
...
*****
calc summary: Total number of verification cases: 3
calc summary: Average call time of all calls (usec): 3.3E-2
*****

```

This shows that the kernel is compiled with success and the correctness check shows that kernel generated data is identical to data generated from original production application. This verification repeated per each of three data files. It also shows that the kernel execution took 3.7E-002 usec for the first verification and 3.3E-2 usec for all three verifications.

3.2 simple-MPI example

In this section, an example of generating a kernel of “calc” from “calc_mod.F90” source file in simple-MPI folder is explained. This example shows a workflow of using KGen for MPI

applications. Especially, this example includes the usage example for using “-e” options which is an important part of supporting large-sized application.

program.F90	update_mod.F90	calc_mod.F90
<pre> PROGRAM calc USE update_mod, only : & update INCLUDE 'mpif.h' INTEGER t, rank, N, err CALL mpi_init(err) CALL mpi_comm_size(...) CALL mpi_comm_rank(...) DO t=1,10 CALL update(rank, N) END DO CALL mpi_finalize(err) END PROGRAM </pre>	<pre> MODULE update_mod USE calc_mod, only : calc PUBLIC update CONTAINS SUBROUTINE update(rank, N) INCLUDE 'mpif.h' INTEGER, INTENT(IN)::rank,N INTEGER :: i, j, err, lsum, gsum(N) INTEGER :: output(COL,ROW) gsum = 0 DO i=1,COL DO j=1,ROW CALL calc(i, j, output) END DO END DO lsum = SUM(output) CALL mpi_gather(lsum, ...) IF (rank==0) THEN print *, 'global sum=', SUM(gsum) END IF END SUBROUTINE END MODULE </pre>	<pre> MODULE calc_mod PUBLIC calc CONTAINS SUBROUTINE calc(i, j, output) INTEGER, INTENT(IN) :: i, j INTEGER, INTENT(OUT), & dimension(:,:) :: output CALL print_msg('start') output(i,j) = i + j CALL print_msg('finish') END SUBROUTINE SUBROUTINE print_msg(msg) CHARACTER(*),INTENT(IN):: & msg PRINT *, msg END SUBROUTINE END MODULE </pre>

In this example callsite is “CALL calc(i, j, output)” in update_mod.F90.

Please change FC and FC_FLAGS in src/Makefile if required to meet your compiler and MPI environment.

To extract “calc” subroutine and make it stand-alone application, following KGen command is used.

NOTE: Following macros are assumed in the command line.

SRC_DIR := \${PWD}/src

SRC := \${SRC_DIR}/update_mod.F90

CALLSITE := update_mod:update:calc

```

>> ${KGen} \
    ${SRC}:${CALLSITE} \
    --cmd-clean "cd ${SRC_DIR}; make -f Makefile.mpirun clean" \
    --cmd-build "cd ${SRC_DIR}; make -f Makefile.mpirun build" \
    --cmd-run "cd ${SRC_DIR}; make -f Makefile.mpirun run" \
    --mpi enable \
    --invocation 0:0:1,0:0:3,1:0:1,1:0:3 \

```



```
--timing repeat=10000 \  
--rebuild all
```

Only “--mpi” flag and “--invocation” flag are explained in this section. Please see a section for “simple” example for the other options.

“--mpi enable”: “--mpi” flag specifies MPI-related information. “enable” sub-options specifies that KGen extracts a kernel from MPI application. Similarly, if target application used OpenMP features, “--openmp enable” option should be added.

“--invocation 0:0:1,0:0:3,1:0:1,1:0:3”: “--invocation” flag specifies when to generate state data in terms of MPI ranks, OpenMP numbers and Nth calls to the kernel region. In this example, four data files will be created from second and fourth invocation from MPI rank 0 and 3. “0” for OpenMp number is used to specify non-OpenMP application such as “1:0:3”

Once KGen completes kernel generation successfully, kernel source files and data files are generated in “kernel” sub-folder as shown below.

kernel folder

```
kernel  
|-- calc.0.0.1  
|-- calc.0.0.3  
|-- calc.1.0.1  
|-- calc.1.0.3  
|-- calc_mod.F90  
|-- kernel_driver.f90  
|-- kgen_statefile.lst  
|-- kgen_utils.f90  
|-- Makefile  
`-- update_mod.F90
```

Four data files, “calc.0.0.1”, “calc.0.0.3”, “calc.3.0.1”, and “calc.3.0.1” in this example, are created in “kernel” folder.

Final step is to build and run the “calc” kernel. Move to “kernel” folder and run “make”. If it completes with success, you will see something similar to following figure on screen.

```

>> make
ifort -c -o kgen_utils.o kgen_utils.f90
ifort -c -o calc_mod.o calc_mod.F90
ifort -c -o update_mod.o update_mod.F90
ifort -c -o kernel_driver.o kernel_driver.f90
ifort -o kernel.exe kernel_driver.o
      update_mod.o kgen_utils.o calc_mod.o

./kernel.exe

***** Verification against 'calc.0.0.1' *****

Number of verified variables:      1
Number of identical variables:    1
Number of non-identical variables within tolerance:      0
Number of non-identical variables out of tolerance:      0
Tolerance:  1.0000000000000000E-014

Verification PASSED

calc : Time per call (usec):  6.000000052154064E-03

***** Verification against 'calc.1.0.1' *****
...
***** Verification against 'calc.1.0.1' *****
...
***** Verification against 'calc.1.0.3' *****
...

*****
calc summary: Total number of verification cases: 4
calc summary: Average call time of all calls (usec):  6.0E-03
*****

```

This shows that the kernel is compiled with success and the correctness check shows that kernel generated data is identical to data generated from original production application. This verification repeated per each of three data files. It also shows that the kernel execution took 6.00E-003 usec for the first verification and 2.19E-002 usec for all four verifications.

3.3 simple-region example

In this section, an example of generating a kernel of “calc” from “calc_mod.F90” source file in simple-region folder is explained. This example shows a workflow of extraction a region of source codes.

program.F90	update_mod.F90	calc_mod.F90
<pre> PROGRAM calc USE update_mod, only : & update INCLUDE 'mpif.h' INTEGER t, rank, N, err CALL mpi_init(err) CALL mpi_comm_size(...) CALL mpi_comm_rank(...) DO t=1,10 CALL update(rank, N) END DO CALL mpi_finalize(err) END PROGRAM </pre>	<pre> MODULE update_mod USE calc_mod, only : calc PUBLIC update CONTAINS SUBROUTINE update(rank, N) INCLUDE 'mpif.h' INTEGER, INTENT(IN)::rank,N INTEGER :: i, j, err, lsum, gsum(N) INTEGER :: output(COL,ROW) gsum = 0 !\$KGEN begin_callsite calc DO i=1,COL DO j=1,ROW CALL calc(i, j, output) END DO END DO !\$KGEN end_callsite lsum = SUM(output) CALL mpi_gather(lsum, ...) IF (rank==0) THEN print *, 'global sum=', SUM(gsum) END IF END SUBROUTINE END MODULE </pre>	<pre> MODULE calc_mod PUBLIC calc CONTAINS SUBROUTINE calc(i, j, output) INTEGER, INTENT(IN) :: i, j INTEGER, INTENT(OUT), & dimension(:, :) :: output CALL print_msg('start') output(i,j) = i + j CALL print_msg('finish') END SUBROUTINE SUBROUTINE print_msg(msg) CHARACTER(*), INTENT(IN):: & msg PRINT *, msg END SUBROUTINE END MODULE </pre>

This example is exactly the same to simple-MPI example except it has “begin_callsite” and “end_callsite” KGen directives in update_mod.F90. Please see simple-MPI example section for the explanations of extracting, building, and executing kernel.

4. Command line options

The syntax of each options generally follows the following convention:

General KGen option syntax:

```
-[-]<option-name> [<suboption-name>=<suboption-value>[, [<suboption-name>=<suboption-value>]]
```

If there are multiple information in <suboption-value>, each information would be separated by colon, “:”. Double or single quotation marks can be used to use some of the separation symbols, such as equal sign, comma, colon, in option value.

[-c, --cmd-clean]

meaning: This mandatory option is Linux command(s) that ensure that next “build commands” actually compiles all source files that KGen requires to analyze.

example) --cmd-clean “cd src; make clean”

`[-b, --cmd-build]`

meaning: This mandatory option is Linux command(s) that compiles target application.

example) `--cmd-build "cd src; make build"`

`[-r, --cmd-run]`

meaning: This mandatory option is Linux command(s) that executes target application.

example) `--cmd-run "cd src; make run"`

`[--outdir]`

meaning : KGen output directory

example) `--outdir /path/to/output/directory`

`[--rebuild]`

meaning : This option forces KGen generates intermittent files such as strace log files and include.ini files. Current version supports "strace", "include", and "state" sub-options. "strace" forces to rebuild strace.log file. "include" forces to rebuild "include.ini" file. "state" forces to rebuild state data files. "All" sub-option is the same to using all of the three sub-options.

example) `--rebuild strace,include,state`

`[--prerun]`

meaning : This options provide a way for user to specify Linux commands that are executed before KGen executes Linux shell command at several stages during kernel extraction. There are five sub-options are supported in this version: "clean", "build", "run", "kernel_build" and "kerne_run". A argument of each sub-commands are executed before executing clean commands, build commands and run command for target application, and build command and run command for KGen generated kernel.

example) `--prerun build="module load intel; module load impi; module load mkl"`

`[--strace]`

meaning : specify paths for strace log file. If starce log file is specified, KGen uses the file instead of generating new strace log file.

example) `--strace ./strace.log`

`[-i, --include-ini]`

meaning: specify paths for include INI file. If include INI file is specified, KGen uses the file instead of generating new include INI file.

example) `--include-ini ./include.ini`

`[--invocation]`

meaning : specifies when to generated state data from which MPI ranks and OpenMP threads.

syntax: `mpi_rank:openmp_num:invocation[,mpi_rank:openmp_num:invocation[...]]`

examples

--invocation 0:1:2 => mpi rank0, openmp num 1, and third invocation of the kernel(starts from 0)

--invocation 1-2:3-4:5-6 => mpi rank1 and 2, openmp num 3 and 4, and sixth and seventh invocations of the kernel

Use 0 for "non MPI application" and use 0 for "non OpenMP application" in the first and second part of the syntax.

[-e, --exclude-ini]

meaning: specify paths for an exclude INI file

example) --exclude ./exclude.ini

[--kernel-option]

meaning : compiler-specific information used in generating Makefile for kernel. Two sub-options are defined in this version: "FC" and "FC_FLAGS". User can choose which Fortran compiler to be used in the kernel makefile with "FC" flag. If user also provide the same information in include.ini file, FC in this option overwrite previous setting and FC_FLAGS in this option added to one in included.ini.

example) --kernel-option FC=ifort,FC_FLAGS=-O3

[--mpi]

meaning : Turns on MPI supports in KGen. There are several sub-options: "enable", "comm", "use", and "header". "enable" specifies that KGen extracts a kernel from MPI application. This is a mandatory for MPI application. "comm" specified the names of variable that is used when MPI call is made. Default "comm" is "MPI_COMM_WORLD". "use" specifies Fortran module name whose name is inserted in additional Fortran use statement. There is no default value for "use". "header" specifies the path to MPI header file. Default "header" is "mpif.h".

example) --mpi ranks=0,comm=mpicom,use="spmd_utils:mpicom"

[--openmp]

meaning : Turns on OpenMP supports in KGen. There is two sub-options: "enable" and "kernel-in-critical-region". "enable" specifies that KGen extracts a kernel from OpenMP application. This is a mandatory for OpenMP application. "kernel-in-critical-region" can has one of two values: "yes" or "no". If "kernel-in-critical-region" is set to "yes", kernel region is encompassed by OpenMP Critical region, which improves correctness of generated state data. However, this may cause deadlock if OpenMP Barrier is used within the kernel.

example) --openmp kernel-in-critical-region=yes

[--intrinsic]

meaning : options to let KGen skip searching for names of intrinsic-procedures. At minimum, one of "skip" or "noskip" should be provided. With "except" sub-flag, user can specify which "namepath" should be considered as exception. With "add-intrinsic" sub-flag, user can add new intrinsic function names. default: --intrinsic skip

example) --intrinsic skip,except=mod_A.subr_B.sum

[--timing]

meaning : provides information about performance measurement. One sub-flag is defined in this version: repeat. "repeat" provides the number of invocations to the kernel subprogram to enhance the measurement resolution.

example) --timing repeat=1000

[--verbose]

meaning: This flag sets the initial verbosity level in the generated kernel. Default value is 1. User can modify the verbosity level by changing the verbosity value that is hard coded in the generated callsite file.

[--check]

meaning: This flag provides KGen with correctness check-related information. Current implementation only allows perturbation related information. "pert_invar" sub-flag select an input variable for perturbation test. "Pert_lim" sub-flag sets the magnitude of perturbation. Default value is '1.0E-15'.

example) --check pert_invar=varname,pert_lim=1.0E-7

[--add-mpi-frame]

meaning: This flag specify to create MPI framework for replicating kernel execution across multiple MPI ranks. This is simple duplication of kernel execution without having any communication among kernels. Two sub-options are allowed: "np" and "mpiexec". "np" sets the number of MPI ranks and "mpiexec" sets the path to mpiexec.

example) --add-mpi-frame=np=4,mpiexec=mpirun

[--source]

meaning : this options specifies information related to source file. "format" sub-flag specifies the Fortran format of source files "fixed" is used for F77 format and "free" used for F90 and later format. With this sub-flag, KGen forces to use the specified format regardless of file extension. "strict" format let parser of KGen inform to apply format strictly or not. Default is of the sub-flag is "no", "alias" sub-flag create path alias. This is useful if you have one file physical location but has two different paths that points the same physical path.

example) format=free,strict=no,alias=/path/A:/path/B

5. Known Issues

- Only subset of Fortran specification is supported.
- A variable of assumed size array is not supported for state generation
- Cyclic linked list is not supported.
- Pointer variable that is associated with part of input state to the kernel may (or may not) generate issues depending on the usage of the variable within the extracted kernel

6. Changes from KGen ver. 0.6.3

6.1 User Interface

- Three mandatory options(clean, build and run of target application) are added in command line.
- strace, rebuild, prerun options are added

6.2 Major Improvements

- Macro definitions and include paths are automatically generated by KGen
- Supports manual state data generation
- Supports OpenMPI